

# HUDM6052 Psychometric II Homework\_04

Chenguang Pan

2023-11-14

## Contents

Q1-a-Parameters-Estimation . . . . .	1
Q1-a-(1) . . . . .	4
Q1-a-(2) . . . . .	4
Q1-a-(3) . . . . .	5

## Q1-a-Parameters-Estimation

*Fit the 1PL, 2PL, and 3PL models...report the estimated item parameters in separated tables*

### My Solution:

To make the layout concise and good-looking, I intentionally omitted the codes for data cleaning and some instant display of running outcomes. I attached the estimated item parameters from all three models into one table to save space.

```
> # load the binary response dataset
> library(mirt)
> df <- read.csv("/Users/panpeter/Desktop/PhD_Learning/HUDM6052 Psychometric II/HUDM6052_Psychometric_II")
> df <- df[,-1]
>
> # -----
> #                               Run 1PL
> # -----
>
> # specify the model for 33 items loading on 1 dimension
> # and constrain all the item slope to be equal for 1PL estimation
> spec <- 'F = 1-33
+ CONSTRAIN = (1-33, a1)'
>
> # estimated the model
> # since I constrained all slopes to be equal, here the argument "2PL" is safe
> irt_1pl <- mirt(df, model = spec, itemtype = "2PL", SE=T)
Iteration: 1, Log-Lik: -20276.377, Max-Change: 0.18304Iteration: 2, Log-Lik: -20265.381, Max-Change: 0.00000
Calculating information matrix...
> irt_1pl
Call:
```

```
mirt(data = df, model = spec, itemtype = "2PL", SE = T)
```

Full-information item factor analysis with 1 factor(s).  
Converged within 1e-04 tolerance after 25 EM iterations.

mirt version: 1.40

M-step optimizer: BFGS

EM acceleration: Ramsay

Number of rectangular quadrature: 61

Latent density type: Gaussian

Information matrix estimated with method: Oakes

Second-order test: model is a possible local maximum

Condition number of information matrix = 50.40657

Log-likelihood = -20260.22

Estimated parameters: 66

AIC = 40588.44

BIC = 40756.74; SABIC = 40648.75

```
> # -----
```

```
> #                               Run 2PL
```

```
> # -----
```

```
> irt_2pl <- mirt(df, model = 1, itemtype = "2PL", SE=T)
```

Iteration: 1, Log-Lik: -20276.377, Max-Change: 0.63072Iteration: 2, Log-Lik: -20095.717, Max-Change: 0.1

Calculating information matrix...

```
> irt_2pl
```

Call:

```
mirt(data = df, model = 1, itemtype = "2PL", SE = T)
```

Full-information item factor analysis with 1 factor(s).  
Converged within 1e-04 tolerance after 32 EM iterations.

mirt version: 1.40

M-step optimizer: BFGS

EM acceleration: Ramsay

Number of rectangular quadrature: 61

Latent density type: Gaussian

Information matrix estimated with method: Oakes

Second-order test: model is a possible local maximum

Condition number of information matrix = 11.16451

Log-likelihood = -20078.62

Estimated parameters: 66

AIC = 40289.23

BIC = 40615.92; SABIC = 40406.3

```
> # -----
```

```
> #                               Run 3PL
```

```
> # -----
```

```
> # specify the model
```

```
> spec <- 'F = 1-33'
```

```

+ PRIOR = (1-33, g, norm, -1.1, 2)'
> irt_3pl <- mirt(df, model = spec, itemtype = "3PL", SE = T)
Iteration: 1, Log-Lik: -20990.083, Max-Change: 1.61517Iteration: 2, Log-Lik: -20300.703, Max-Change: 1.4

```

Items	1PL				2PL				3PL			
	a	b	g	u	a	b	g	u	a	b	g	u
mathc1	0.960	-0.881	0	1	1.078	-0.816	0	1	1.385	-0.250	0.238	1
mathc2	0.960	-1.706	0	1	1.135	-1.516	0	1	1.198	-1.170	0.198	1
mathc3	0.960	-0.162	0	1	1.307	-0.145	0	1	1.529	0.080	0.095	1
mathc4	0.960	-0.412	0	1	1.308	-0.349	0	1	2.095	0.170	0.235	1
mathc5	0.960	1.104	0	1	0.627	1.556	0	1	2.282	1.443	0.187	1
mathc6	0.960	-0.696	0	1	1.352	-0.569	0	1	1.428	-0.432	0.060	1
mathc7	0.960	-0.457	0	1	1.044	-0.434	0	1	1.247	-0.081	0.141	1
mathc8	0.960	-0.562	0	1	0.897	-0.589	0	1	0.951	-0.395	0.071	1
mathc9	0.960	-2.089	0	1	1.415	-1.609	0	1	1.464	-1.443	0.129	1
mathc11	0.960	-0.668	0	1	1.243	-0.571	0	1	1.956	0.092	0.293	1
mathc12	0.960	0.470	0	1	0.856	0.514	0	1	1.130	0.802	0.114	1
mathc13	0.960	-1.212	0	1	1.535	-0.910	0	1	2.546	-0.207	0.360	1
mathc14	0.960	0.552	0	1	0.775	0.655	0	1	0.887	0.819	0.060	1
mathc15	0.960	0.498	0	1	0.726	0.624	0	1	0.858	0.838	0.075	1
mathc16	0.960	-0.216	0	1	0.829	-0.237	0	1	0.942	0.027	0.090	1
mathc17	0.960	0.715	0	1	0.837	0.795	0	1	1.254	1.050	0.125	1
mathc18	0.960	0.767	0	1	1.107	0.688	0	1	1.253	0.757	0.034	1
mathc19	0.960	0.290	0	1	0.551	0.475	0	1	1.432	1.302	0.297	1
mathc21	0.960	-0.004	0	1	0.607	0.011	0	1	3.411	1.149	0.407	1
mathc22	0.960	-0.253	0	1	0.990	-0.250	0	1	1.401	0.295	0.212	1
mathc23	0.960	-0.085	0	1	0.688	-0.099	0	1	1.428	0.863	0.316	1
mathc24	0.960	-0.052	0	1	0.986	-0.053	0	1	1.460	0.460	0.202	1
mathc25	0.960	-0.170	0	1	1.230	-0.155	0	1	1.517	0.132	0.122	1
mathc26	0.960	0.332	0	1	0.893	0.350	0	1	1.177	0.672	0.125	1
mathc27	0.960	-0.282	0	1	1.479	-0.232	0	1	1.719	-0.018	0.094	1
mathc28	0.960	-0.026	0	1	1.949	-0.044	0	1	2.260	0.100	0.058	1
mathc29	0.960	0.119	0	1	1.229	0.092	0	1	1.928	0.476	0.173	1
mathc31	0.960	1.250	0	1	0.723	1.564	0	1	1.699	1.576	0.150	1
mathc32	0.960	-0.476	0	1	0.837	-0.522	0	1	1.600	0.481	0.350	1
mathc33	0.960	0.176	0	1	1.284	0.134	0	1	2.790	0.582	0.217	1
mathc34	0.960	0.165	0	1	0.678	0.229	0	1	0.863	0.645	0.131	1
mathc35	0.960	0.667	0	1	0.360	1.577	0	1	1.480	1.899	0.288	1
mathc36	0.960	1.449	0	1	1.060	1.352	0	1	1.601	1.359	0.070	1

Due to the limitation of `mirt` package, I can't constrain all the  $\alpha$  to be 1. Rather, I can only set them to be equal across all the items. Therefore, in the estimation for the 1PL, the estimated universal  $\alpha$  is .96 here.

### Q1-a-(1)

*Does it appear reasonable to assume all the items having an equal slope...*

#### My Solution:

No.

From a aspect of test development, since this is a test about math placement, we should expect that items can discriminate students with different traits well. In addition, comparing the estimated parameters from 2PL model versus 1PL, these items' levels of discrimination spread along a wide range. It is reasonable to have items with higher levels of discrimination than others.

From a mathematics perspective, since the 1PL model is nested in the 2PL model, I conducted the Likelihood Ratio Test to compare the two models as followed:

$$D = -2[\ln(L_{1pl}) - \ln(L_{2pl})].$$

Plug the log likelihood estimated from the above code chunk, then one can have  $D = 363.2$  at the degree of freedom of  $df = df_{2pl} - df_{1pl} = 66 - 34 = 32$ . Based on the Chi-squared distribution, the p value is lower than .001. Therefore, 2PL is better than 1PL, which means the discrimination is preferred.

### Q1-a-(2)

*Does it appear useful to include a guessing parameter in the model...*

#### My Solution:

Yes, it is useful.

Intuitively, it is reasonable to include a guessing parameter since this a test with multiple choice and guessing is very possible. In addition, by looking through all the guessing parameters, one can find that the `matchc13`, `matchc21`, and `matchc32` do have quite high guessing rate, i.e., all above .30.

However, in terms of model comparison, when using the LRT test again to compare the 2PL vs 3PL model, one can have  $D = 28.28$  at 33 degree of freedom,  $P = .701$ . Based on the parsimony rule, one should endorse the simpler model, i.e., the 2PL.

Therefore, my overall conclusion is including a guessing parameter is useful in this scenario. A practitioner should choose the either model based on their purpose since these two models do not differ a lot.

```
> # get the fit indices for 1PL model
> M2(irt_1pl, na.rm = T)
      M2 df p      RMSEA      RMSEA_5      RMSEA_95      SRMSR      TLI
stats 992.2492 527 0 0.03104487 0.02805368 0.03398001 0.05781072 0.9632371
      CFI
stats 0.9633067
>
> # get the fit indices for 2PL model
> M2(irt_2pl, na.rm = T)
      M2 df      p      RMSEA      RMSEA_5      RMSEA_95      SRMSR
stats 678.0246 495 7.852727e-08 0.02009113 0.01617592 0.02371771 0.03274386
      TLI      CFI
stats 0.9846029 0.9855652
```

```

>
> # get the fit indices for 3PL model
> M2(irt_3pl, na.rm = T)
      M2    df          p      RMSEA    RMSEA_5    RMSEA_95      SRMSR
stats 604.7095 462 8.359281e-06 0.01836359 0.01400902 0.02228458 0.03106851
      TLI      CFI
stats 0.9871369 0.9887448

```

### Q1-a-(3)

Evaluate the goodness of fit of the items with the option of the chi-square test...

#### My Solution:

I conduct the item fit analysis on each model and summarize the results into one table to make the layout concise.

```

> # get the item fit indices for each model
> item_fit_1pl <- itemfit(irt_1pl, na.rm = T)
> item_fit_2pl <- itemfit(irt_2pl, na.rm = T)
> item_fit_3pl <- itemfit(irt_3pl, na.rm = T)
>
> # combine all the outputs into one table
> item_fit_all <- cbind(item_fit_1pl[,c("item")],
+                       round(item_fit_1pl[,c("S_X2", "p.S_X2")],4),
+                       round(item_fit_2pl[,c("S_X2", "p.S_X2")],4),
+                       round(item_fit_3pl[,c("S_X2", "p.S_X2")],4))
> names(item_fit_all)[1] <- "item"
>
> # get all the item fit indices for 1PL, 2PL, and 3PL model
> item_fit_all
      item      S_X2 p.S_X2      S_X2 p.S_X2      S_X2 p.S_X2
1  mathc1 24.8711 0.4128 23.0566 0.4575 23.8800 0.3536
2  mathc2 28.1234 0.2112 30.4743 0.1074 28.6024 0.1239
3  mathc3 27.2722 0.2919 22.4566 0.4929 21.7064 0.4775
4  mathc4 26.5449 0.3261 19.2254 0.6314 17.4558 0.6831
5  mathc5 97.3404 0.0000 52.8362 0.0014 39.3238 0.0183
6  mathc6 35.2844 0.0643 26.8461 0.2171 23.9949 0.2933
7  mathc7 30.6375 0.1645 30.4324 0.1708 30.8269 0.1271
8  mathc8 31.1674 0.1490 31.5235 0.1393 23.0747 0.4564
9  mathc9 31.5513 0.0854 16.4299 0.6896 12.8670 0.7994
10 mathc11 24.7300 0.4205 17.4787 0.7364 16.8747 0.7187
11 mathc12 23.5755 0.4861 22.8300 0.5298 20.6982 0.5995
12 mathc13 35.9083 0.0422 24.6484 0.2627 19.3900 0.3682
13 mathc14 27.5382 0.2800 23.1188 0.5128 20.5069 0.6112
14 mathc15 33.5261 0.0934 21.0409 0.6903 17.2351 0.7976
15 mathc16 26.6813 0.3195 25.3271 0.3882 22.7550 0.5343
16 mathc17 21.2673 0.6229 17.9679 0.8046 17.1584 0.8014
17 mathc18 21.5747 0.6046 22.3173 0.5012 19.1247 0.6376
18 mathc19 54.0475 0.0004 20.7629 0.7058 20.1995 0.6854
19 mathc21 61.7627 0.0000 33.2577 0.1247 23.3450 0.4408
20 mathc22 20.4045 0.6736 20.4046 0.6736 20.8114 0.5926
21 mathc23 20.8365 0.6483 14.4939 0.9524 13.5269 0.9566
22 mathc24 20.9306 0.6428 21.0813 0.6339 21.0628 0.5773

```

23	mathc25	24.9448	0.4088	21.3330	0.5608	21.4360	0.4939
24	mathc26	19.3543	0.7328	17.6730	0.8186	18.2240	0.7452
25	mathc27	45.0529	0.0057	29.1187	0.1112	27.4429	0.1567
26	mathc28	65.4501	0.0000	24.1670	0.2352	23.9470	0.1982
27	mathc29	20.6986	0.6564	16.3725	0.8389	14.2529	0.8923
28	mathc31	38.2013	0.0242	24.7382	0.4771	22.2939	0.5026
29	mathc32	26.7508	0.3162	22.5725	0.5451	21.6350	0.5424
30	mathc33	29.0666	0.2176	21.2045	0.5686	8.0644	0.9949
31	mathc34	37.3521	0.0403	26.8718	0.3623	24.8798	0.4123
32	mathc35	109.6358	0.0000	32.7703	0.2048	35.3966	0.1033
33	mathc36	28.9987	0.2202	25.1823	0.3410	23.7681	0.3595

The results show that all items fit well in all models except the **mathc5** item.